

- Managing SELinux Security

- Implementing Advanced Storage Features

-----  
**Managing SELinux Security**

**CHANGING THE SELINUX ENFORCEMENT MODE**

we have 2 security models:

1-Discretionary Access Control-DAC

- permission
- special permission
- adv. permission
- sudo

2-Mandatory Access Control-MAC

- SELinux

**SELinux,**

- its set of security rules that determine which process can access to which file, directory and ports.
- every file, dir, process and port they have special security label called SELinux context.
- A context is a name used by SELinux policy to determine whether a process can access to file, dir and port.
- by default SELinux policy won't allow any interaction unless explicitly rule granted access.
- SELinux protects user's data by attaching security context label to file, dir, port, process
- SELinux security label has several parts:

**user**

**role**

**type**

**sensitivity**

between them **type** is important.

**SELinux modes:**

- 1-enforcing** SELinux security policy is enforced.
- 2-permissive** SELinux prints warnings instead of enforcing.
- 3-disabled** No SELinux policy is loaded.

**current SELinux mode**

# getenforce

Enforcing

# sestatus

**change SELinux mode:**

**temporary**

# setenforce 1/0

usage: setenforce [ Enforcing | Permissive | 1 | 0 ]

**persistently**

# vim /etc/selinux/config ->original

SELINUX=**enforcing**

:wq!

or

# vim /etc/sysconfig/selinux ->softlink

**watch SELinux security context**

# touch /tmp/file1

# mkdir /tmp/testdir

**watch DAC**

# ls -l /tmp/file1

-rw-r--r--. 1 root root 0 Jun 12 10:54 /tmp/file1

**watch MAC**

# ls -lZ /tmp/file1

-rw-r--r--. 1 root root **unconfined\_u:object\_r:user\_tmp\_t:s0** 0 Jun 12 10:54 /tmp/file1

## SELinux scenario:

- SELinux on port
- SELinux on file/dir

### -SELinux on port

servera.lab.example.com 172.25.250.10

serverb.lab.example.com 172.25.250.11

#### order:

change ssh port on servera **22/tcp** to **22000/tcp** and make ssh from serverb to servera with new port

#### servera

```
# netstat -lnptu
```

```
tcp      0      0 0.0.0.0:22
```

```
tcp6     0      0 :::22
```

```
# systemctl status sshd.service
```

```
# vim /etc/ssh/sshd_config
```

```
17 Port 22000
```

```
:wq!
```

```
# firewall-cmd --permanent --add-port=22000/tcp
```

```
# firewall-cmd --reload
```

```
# firewall-cmd --list-all
```

```
# systemctl restart sshd.service
```

Job for sshd.service failed because the control process exited with error code.

```
# journalctl | grep "ssh"
```

SELinux is preventing sshd from name\_bind access on the tcp\_socket port 22000.

#### SELinux log

```
# sealert -l 7786ad15-3f9b-42f2-86f1-a1ccd7ff71ef
```

```
# getenforce
```

#### Enforcing

#### Solution

add port 22000 in to SELinux policy database for ssh\_port\_t

```
# semanage port -l
```

```
# semanage port -l | grep "ssh"
```

```
ssh_port_t      tcp      22
```

```
# semanage port -a -t ssh_port_t -p tcp 22000
```

```
-a add
```

```
-d delete
```

```
-t type of security context
```

```
-p protocol tcp/udp
```

```
# systemctl restart sshd.service
```

```
# netstat -lnptu
```

```
tcp      0      0 0.0.0.0:22000
```

```
tcp6     0      0 :::22000
```

#### verify from serverb

```
# ssh root@servera.lab.example.com -p 22000
```

## -SELinux on file/dir

```
# yum search apache
httpd.x86_64 : Apache HTTP Server
# yum info httpd.x86_64
# yum list httpd.x86_64
# yum install httpd.x86_64 -y
# systemctl list-unit-files | grep "http"
httpd.service                                disabled
# systemctl enable httpd.service
# systemctl start httpd.service
# systemctl status httpd.service
# firewall-cmd --permanent --add-port=80/tcp
# firewall-cmd --reload
# firewall-cmd --list-all
-DocumentRoot ->/var/www/html/
-DirectoryIndex ->index.html
DocumentRoot/DirectoryIndex ->/var/www/html/index.html
by default, when Install Apache web-server DocumentRoot will create but DirectoryIndex should be
# echo "Hello Class!" >/var/www/html/index.html
# cat /var/www/html/index.html
Hello Class!
# systemctl restart httpd.service
# curl localhost
Hello Class!
```

## check MAC

```
# ls -ldZ /var/
drwxr-xr-x. 22 root root system_u:object_r:var_t:s0 4096 Jun 12 11:18 /var/
# ls -ldZ /var/www/
drwxr-xr-x. 4 root root system_u:object_r:httpd_sys_content_t:s0 33 Jun 12 11:18 /var/www/
# ls -ldZ /var/www/html/
drwxr-xr-x. 2 root root system_u:object_r:httpd_sys_content_t:s0 24 Jun 12 11:27 /var/www/html/
# ls -ldZ /var/www/html/index.html
-rw-r--r--. 1 root root unconfined_u:object_r:httpd_sys_content_t:s0 13 Jun 12 11:27 /var/www/html/index.html
```

all got inherit from parent

-index.html from html dir

-html dir from www dir

```
# semanage fcontext -l
```

```
# semanage fcontext -l | grep "/var/www(/.*)?"
```

```
/var/www(/.*)?                all files          system_u:object_r:httpd_sys_content_t:s0
```

```
(/.*)? ->directory + whatever will create inside it.
```

## order:

change DocumentRoot to /virtual

```
# mkdir /virtual
```

```
# echo "Helloooooo!" >/virtual/index.html
```

```
# ls -ldZ /virtual/
```

```
drwxr-xr-x. 2 root root unconfined_u:object_r:default_t:s0 24 Jun 12 11:36 /virtual/
```

```
# ls -lZ /virtual/index.html
```

```
-rw-r--r--. 1 root root unconfined_u:object_r:default_t:s0 12 Jun 12 11:36 /virtual/index.html
```

```
# vim /etc/httpd/conf/httpd.conf
```

```
122 DocumentRoot "/virtual"
```

```
127 <Directory "/virtual">
```

```
135 <Directory "/virtual">
```

```
:wq!
```

```
# systemctl restart httpd.service
```

```
# curl localhost
```

it shows you:

## HTTP SERVER TEST PAGE

but me expect is:

**Helloooooo!**

because SELinux stopped access to /virtual/index.html

```
# journalctl | grep "http"
```

SELinux is preventing httpd from getattr access on the file /virtual/index.html.

### SELinux log

```
# sealert -l 10f00d67-26b2-4a26-b025-d5f91d82831f
```

SELinux is preventing httpd from getattr access on the file /virtual/index.html.

If you want to allow httpd to have getattr access on the index.html file

Then you need to change the label on /virtual/index.html

Do

```
# semanage fcontext -a -t FILE_TYPE '/virtual/index.html'
```

### Solution

change /virtual directory and whatever inside it available context to apache acceptable context for SELinux.

#### 1-temporary:

```
# chcon -Rv -t httpd_sys_content_t /virtual/
```

changing security context of '/virtual/index.html'

changing security context of '/virtual/'

```
# ls -ldZ /virtual/
```

```
drwxr-xr-x. 2 root root unconfined_u:object_r:httpd_sys_content_t:s0 24 Jun 12 11:36 /virtual/
```

```
# ls -lZ /virtual/index.html
```

```
-rw-r--r--. 1 root root unconfined_u:object_r:httpd_sys_content_t:s0 12 Jun 12 11:36 /virtual/index.html
```

```
# curl localhost
```

Helloooooo!

**restore** to default SELinux security context

```
# restorecon -Rv /virtual/
```

Relabeled /virtual from unconfined\_u:object\_r:httpd\_sys\_content\_t:s0 to unconfined\_u:object\_r:default\_t:s0

Relabeled /virtual/index.html from unconfined\_u:object\_r:httpd\_sys\_content\_t:s0 to unconfined\_u:object\_r:default\_t:s0

```
# curl localhost
```

it shows you:

**HTTP SERVER TEST PAGE**

#### 2-persitently

change /virtual security context in to SELinux policy database from current to what should be:

```
# semanage fcontext -a -t httpd_sys_content_t "/virtual(/.*)"?"
```

```
# restorecon -Rv /virtual/
```

Relabeled /virtual from unconfined\_u:object\_r:default\_t:s0 to unconfined\_u:object\_r:httpd\_sys\_content\_t:s0

Relabeled /virtual/index.html from unconfined\_u:object\_r:default\_t:s0 to unconfined\_u:object\_r:httpd\_sys\_content\_t:s0

```
# curl localhost
```

Helloooooo!

## Implementing Advanced Storage Features

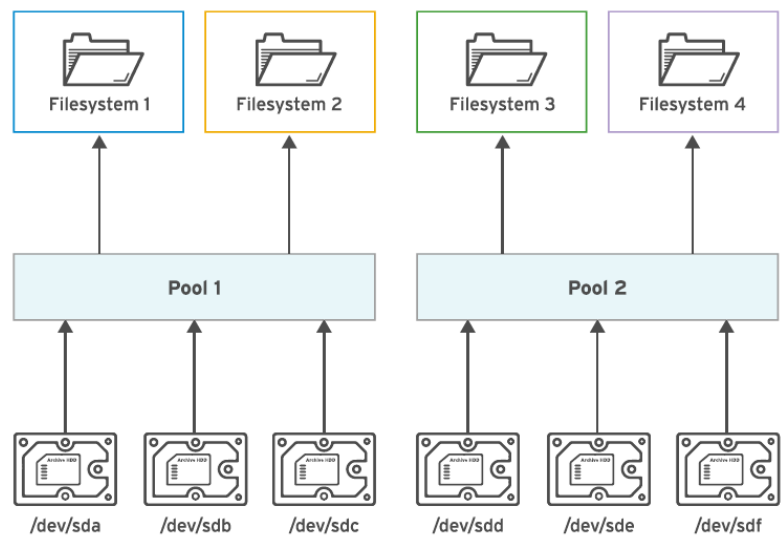
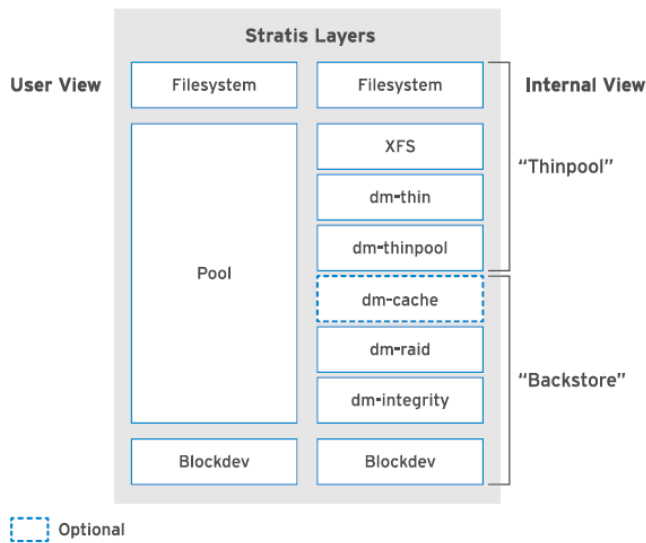
### MANAGING LAYERED STORAGE WITH STRATIS

#### DESCRIBING THE STRATIS ARCHITECTURE

-The current local storage solution in Red Hat Enterprise Linux (RHEL) includes many stable and mature technologies, including the device mapper (dm), the logical volume manager (LVM) and the XFS file system.

-With RHEL8, Red Hat introduces the Stratis local storage management solution.

Instead of developing from scratch, as other storage projects attempted, Stratis works with existing RHEL storage components.



#### Implement Stratis on Linux

##### stratis info

**package:** stratis-cli.noarch, stratisd.x86\_64

**daemon:** stratid.service

**working dir:** /dev/stratis

##### steps:

- create pool
- create filesystem
- mount and use it
- create snapshot

##### Install

```
# yum install stratis* -y
# systemctl enable stratid.service
# systemctl start stratid.service
# systemctl status stratid.service
# lsblk
sdb      8:16  0   5G  0 disk
sdc      8:32  0   5G  0 disk
sdd      8:48  0   5G  0 disk
# stratis
blockdev daemon filesystem -h --help key pool --propagate report --version
```

##### -create pool

```
# stratis pool create <pool name> <storage name>
```

```
# stratis pool create pool1 /dev/sdb
```

```
# stratis pool list
```

```
# stratis blockdev list
```

##### append storage to existing pool

```
# stratis pool add-data <existing pool> <new storage>
```

```
# stratis pool add-data pool1 /dev/sdc
```

```
# stratis pool list
```

```
# stratis blockdev list
```

##### -create filesystem

```
# stratis filesystem create <existing pool> <file-system name>
```

```
# stratis filesystem create pool1 fs1
```

```
# stratis filesystem list
```

```
# blkid
```

```
/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-43779909bd5e4af7a071f3d4ae4707aa: UUID="43779909-bd5e-4af7-a071-f3d4ae4707aa" TYPE="xfs"
```

```
# tree /dev/stratis/
```

```
/dev/stratis/
```

```
└─ pool1
```

```
    └─ fs1 -> ../dm-7
```

## -mount and use it

```
# mkdir /mnt/disk1
```

```
# echo "/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-43779909bd5e4af7a071f3d4ae4707aa /mnt/disk1 xfs defaults,x-systemd.requires=stratisd.service 0 0" >>/etc/fstab
```

```
# mount -a
```

```
# df -hT
```

```
/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-43779909bd5e4af7a071f3d4ae4707aa xfs 1.0T 7.2G 1017G 1% /mnt/disk1
```

## -create snapshot

```
# stratis filesystem snapshot <existing pool> <existing filesystem> <snapshot name>
```

```
# stratis filesystem snapshot pool1 fs1 fs1-snap
```

```
# stratis filesystem list
```

## restore snapshot

```
# umount /mnt/disk1
```

```
# vim /etc/fstab
```

```
delete fs1 record
```

```
:wq!
```

```
# mount -a
```

```
# stratis filesystem destroy pool1 fs1
```

```
# stratis filesystem list
```

```
# blkid
```

```
/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-81293b8746fb47e4b0350db88a727158: UUID="81293b87-46fb-47e4-b035-0db88a727158" TYPE="xfs"
```

```
# echo "/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-81293b8746fb47e4b0350db88a727158 /mnt/disk1 xfs defaults,x-systemd.requires=stratisd.service 0 0" >>/etc/fstab
```

```
# mount -a
```

```
# df -hT
```

```
/dev/mapper/stratis-1-52a745e6ba094c3eafcd54c57b988238-thin-fs-81293b8746fb47e4b0350db88a727158 xfs 1.0T 7.2G 1017G 1% /mnt/dis
```

```
# ls /mnt/disk1
```

```
file1 shadow
```

## COMPRESSING AND DEDUPLICATING STORAGE WITH VDO

### DESCRIBING VIRTUAL DATA OPTIMIZER

-RHEL8 includes the **Virtual Data Optimizer-VDO** driver, which optimizes the data footprint on block devices.

-VDO is a Linux device mapper driver that reduces disk space usage on block devices, and minimizes the replication of data, saving disk space and even increasing data throughput.

-VDO includes two kernel modules:

the **kvdo** module to transparently control data compression

the **uds** module for deduplication.

### VDO How it works

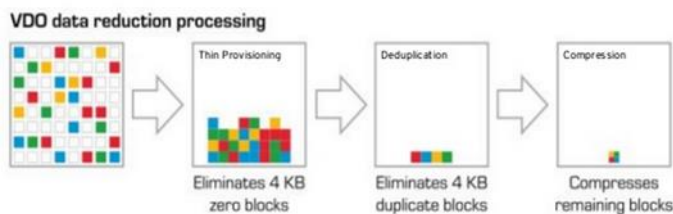
VDO applies three phases to data in the following order to reduce the footprint on storage devices:

**1. Zero-Block Elimination** filters out data blocks that contain only zeroes (0) and records the information of those blocks only in the metadata. The nonzero data blocks are then passed to the next phase of processing. This phase enables the thin provisioning feature in the VDO devices.

**2. Deduplication eliminates redundant data blocks.** When you create multiple copies of the same data, VDO detects the duplicate data blocks and updates the metadata to use those duplicate blocks as references to the original data block without creating redundant data blocks.

The **universal deduplication service-UDS** kernel module checks redundancy of the data through the metadata it maintains.

**3. Compression is the last phase.** The kvdo kernel module compresses the data blocks using LZ4 compression and groups them on 4 KB blocks.



### Implement VDO on Linux

**vdo info**

**package:** vdo.x86\_64, kmod.x86\_64

**daemon:** vdo.service

### Install

```
# yum install kmod.x86_64 vdo.x86_64 -y
```

```
# systemctl enable vdo.service
```

```
# systemctl start vdo.service
```

### create

```
# vdo create --name vdo1 --device /dev/sdd --vdoLogicalSize 50G
```

VDO instance 0 volume is ready at **/dev/mapper/vdo1**

```
# mkfs.xfs -K /dev/mapper/vdo1
```

```
# mkdir /mnt/disk2
```

```
# blkid
/dev/mapper/vdo1: UUID="d147065b-6c72-4470-8c21-faf569b1a68a" TYPE="xfs"
# echo "/dev/mapper/vdo1 /mnt/disk2 xfs defaults,x-systemd.requires=dvo.service 0 0" >>/etc/fstab
# mount -a
# df -hT
/dev/mapper/vdo1      xfs      50G 390M  50G  1% /mnt/disk2
# vdo status --name vdo1
# vdo status --name vdo1 | grep -i "deduplication"
# vdo status --name vdo1 | grep -i "compression"
```